

- Imported the following libraries:
 - NumPy, torch (for Tensors),
 - torch.nn (Neural Network layers),
 - torch.optim (optimizers and improving model parameters),
 - torch.utils.data (for data handling),
 - torchmetrics (for measuring model accuracy),
 - torchvision.datasets (for importing popular vision datasets), and
 - torchvision.transforms (for image processing and data augmentation).
- Built a single callable using `tensor.Compose([...])` by chaining several image transforms in order. Done for the **purpose of data augmentation and changing PIL images to tensors**.
- Applied `tensor.Compose([...])` to both train and test subsets of the FashionMNIST dataset, a de-factor dataset that contains 10 categories of clothing.
- Identified several other variables before proceeding:
 - `Classes = train_data.classes`
 - `Num_classes = len(classes)`
 - `num_input_channels = 1`: As FashionMNIST images are greyscale.
 - `Num_output_channels = 16`: Model Choice, number of feature maps (i.e. outputs of convolutional operator between feature detector and original images) each convolutional layer will produce
 - `Image_size = train_data[0][0].shape[1]`: to extract **height and width** value of (any) training sample/image.
- Declared a PyTorch model class which inherits from **nn.Module** (the base class for all neural network models).
 - Within this class, I implemented the **constructor**, which is a special method in Python automatically called when a new object of the class is created. This constructor defines the architecture of the model by:
 - Receiving `num_classes` as an input argument, which determines how many output scores (logits) the model will produce (e.g., 10 for Fashion-MNIST).
 - Calling `super().__init__()`, which ensures that all internal machinery of `nn.Module` is properly initialized, including parameter tracking, registration of layers, and enabling save/load functionality.
 - Declaring and initializing all layers (convolutional, activation, pooling, flattening, and fully connected) that will be used during the forward pass.
- **Defined the forward pass** (forward method), which specifies how input data is transformed step by step through the network:
 - Passes input images through convolutional layers to extract hierarchical spatial features.
 - Applies activation functions (e.g., ReLU) to introduce non-linearity and enable the model to learn complex patterns and to avoid overfitting.
 - Uses pooling layers to downsample feature maps, reducing computational complexity while retaining essential information.
 - Flattens the high-level feature maps into a vector suitable for classification
 - Processes the flattened vector through fully connected layers to produce class logits corresponding to `num_classes`.
- **Engineered an efficient data handling pipeline**

- Utilized DataLoader to batch and shuffle training samples, improving training throughout and model generalization.
- Ensured compatibility between dataset structure and model input through proper preprocessing and dimensional alignment.
- **Implemented a robust training loop**
 - Defined CrossEntropyLoss as the objective function for multi-class classification tasks.
 - Applied the Adam optimizer with carefully chosen hyperparameters (e.g., learning rate = 0.001) for stable and fast convergence.
 - Automated the forward pass, loss computation, gradient backpropagation, and parameter updates across multiple epochs.
 - Integrated runtime loss tracking (running_loss) to monitor learning progress and detect under/overfitting trends.
- **Optimized for maintainability and experimentation**
 - Encapsulated the entire training procedure in a reusable train_model function to simplify future model iterations.
 - Facilitated model scaling by parameterizing class count (num_classes) and epoch control (num_epochs)
- **Built a performance evaluation pipeline for trained deep learning models**
 - Configured a DataLoader for the test dataset with controlled batching (batch size = 10) and sequential sampling (shuffle=False) to ensure consistent evaluation order.
 - Switched the model to evaluation mode (net.eval()) to deactivate dropout/batch norm updates and guarantee stable inference.
- **Integrated class-wise and overall performance metrics**
 - Implemented accuracy measurement using Accuracy(task='multiclass') to assess global classification correctness.
 - Added Precision and Recall metrics with class-wise granularity (average=None) to diagnose per-class strengths and weaknesses.
- **Automated batched inference and prediction aggregation**
 - Performed forward passes on each batch with input reshaping to match model requirements (features.reshape(-1, 1, image_size, image_size)).
 - Extracted class predictions using torch.argmax, extending results into a consolidated prediction list for later analysis.
- **Computed and reported final evaluation statistics**
 - Aggregated computed metrics via .compute() calls, converting them into Python-native types for reporting and further processing.
 - Produced summary outputs including overall accuracy and per-class precision/recall, enabling interpretable model performance benchmarking.